

# Threads

# Objectives

- Understand the concept of multithreading
- Understand the life cycle of a thread
- Understand thread synchronization
- Be familiar with thread scheduling and prioritization techniques

# Threads

- **Provide**
  - Concurrent paths of execution
  - A way to enable of parallel processing
    - Multi processor/core offer true parallelism
    - Single processor/core give the illusion
- **Why Use Threads?**
  - Improve program performance
  - Execute separate but coordinated tasks
  - Run background tasks

# Java Programs are Multi-Threaded

- Every uses multiple thread
  - The “main” thread executes `main()`
  - Another thread performs garbage collection
- Many library classes use threading techniques to improve perceived performance and protect data integrity

# java.lang.Thread

- `start()` – **housekeeping, becomes runnable**
- `run()` – **where your code is executed**
- `interrupt()` – **sets interruption status**
  - **If blocked on `wait`, `sleep` or `join`**  
**`InterruptedException` is thrown and the status cleared**
- `isInterrupted()` – **returns interruption status**
- `interrupted()` – **returns interruption status of current thread and resets the status**
- `join()` – **waits till the thread terminates**
- `sleep()` – **sleeps for some time**
- `yield()` – **pause and let another execute**
- `set(Default)UncaughtExceptionHandler()` – **installs a handler for uncaught exception**

# java.lang.Thread

- `resume()` - resumes a suspended thread
- `stop()` - stops the thread
- `suspend()` - suspends the thread

# How Threads Run

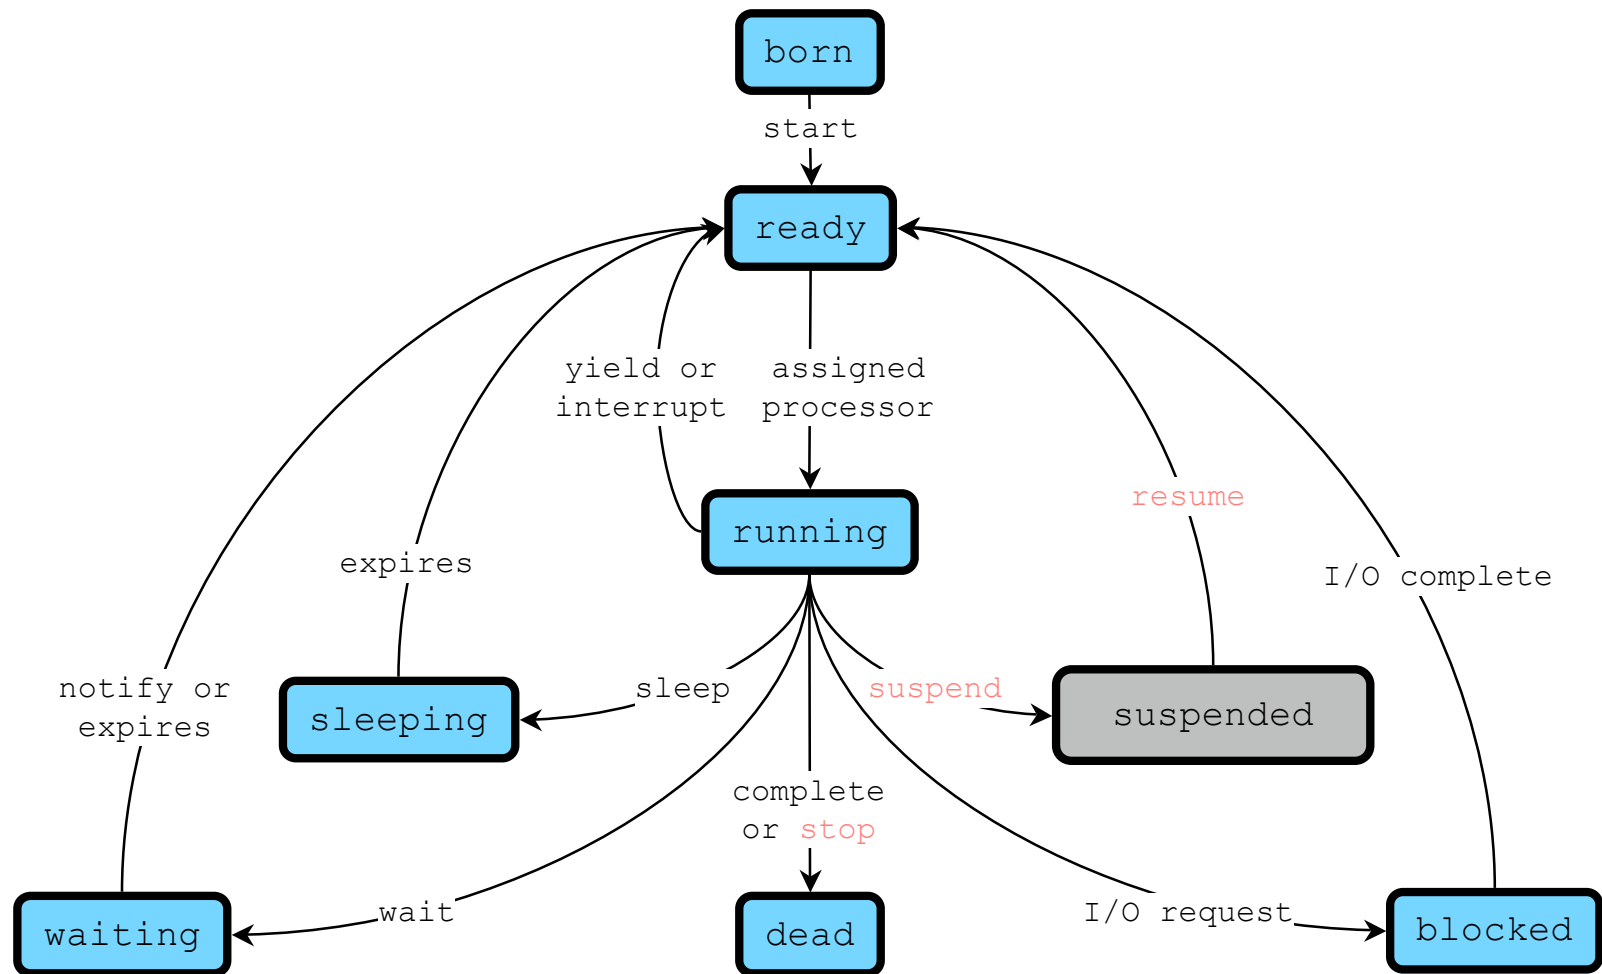
- **Each thread needs time to execute**
  - **Pre-emptive time-slicing**
    - Performed at process level
  - **Scheduling**
    - No scheduling guarantees
    - No fairness guarantees
    - No guarantee threads will make progress

# Thread Priority

- **Only thread scheduler uses priority**
  - **Priority heuristically influences scheduler**
    - 10 Crisis management
    - 7-9 Interactive, event driven
    - 4-6 IO - bound
    - 2-3 Background computation
    - 1 Run only if nothing else can
- `setPriority()`
  - `MIN_PRIORITY, 1`
  - `MAX_PRIORITY, 10`
  - `NORM_PRIORITY, 5`



# Thread Lifecycle



# Controlling Thread State

- Use flag to indicate if thread is running
- Don't be selfish
- Do not use
  - `stop()`
  - `suspend()`
  - `resume()`
- **Implement stop using flags, `interrupt()`, `isInterrupted()` and `interrupted()`**
- **Name threads (useful for debugging)**

# Controlling Thread State

```
volatile boolean running = true;

public void run()
{
    while( running && !Thread.interrupted() )
    {
        // run and do something in a loop
    }
}

public void stopThisThread()
{
    running = false;
    interrupt();
}
```

# Non-portable Thread Features

- The use of priorities and `yield()` to “tune” concurrency should be avoided
- Behavior of `yield()` and impact of priorities vary from VM to VM.

# java.lang.Runnable

- **Interface specifying the `run()` operation**
  - **Implemented by `Thread`**
- **Using `Runnable`**
  - **Do not need to extend from `Thread`**
  - **Separates logic from concurrency**
- **Implementation**
  1. **Define implementing class**
  2. **Provide an instance to the `Thread` constructor**
  3. **Start the thread**

# Synchronization

# Reasons for Synchronization

- **Example**

```
int totalAmount = price * shares;  
// 1. report total amount to the user  
// 2. get money equal to total amount  
// 3. buy shares
```

***Price can be changed by other threads in  
between operations***

***Synchronize access to price***

# Monitor and Synchronization

- **Monitor (one per object)**
  - Implemented in `Object`
  - Only one thread may have ownership of objects monitor at a time
- **Synchronization**
  - An objects monitor is used to restrict access to guarded statements
  - `synchronized` **keyword**
    - Method modifier
    - Controls access to a block



# synchronized

- **A thread may not execute synchronized statements until it obtains ownership of controlling monitor**
  - **Class methods**
    - `Class` object for the method's class
  - **Instance methods**
    - The object the method was invoked on
  - **Synchronized block**
    - The object specified in the synchronized block

# synchronized Block

```
Object m_Lock = new Object();  
public void doSomething()  
{  
    synchronized(m_Lock)  
    {  
        // current thread now has ownership of m_Lock  
    }  
    // current thread releases ownership of m_Lock  
    ...  
}
```

# Atomic Assignment

- Assignment cannot be interrupted

`foo = 50;`

`foo = 42;`

- Any thread will see either 42 or 50, not garbage
- No guarantee value is result of most recent assignment
  - `volatile` – value is “flushed” immediately after written

- Except... long and double

- Unless `volatile`

# Assignment & Synchronization

```
private Object mFoo;  
...  
public void set(Object foo)  
{  
    mFoo = foo;  
}
```

- **Doesn't need to be synchronized**

# Assignment & Synchronization

```
private Object mFoo;  
private boolean mFooAssigned;  
...  
public synchronized void set(Object foo)  
{  
    mFoo = foo;  
    mFooAssigned = true;  
    Syetem.out.println( foo );  
}
```

- **Needs to be synchronized, because:**
  - **Multiple related variables**
  - **Multiple operations**

# Assignment & Synchronization

```
public void set(Object foo)
{
    mFoo = foo;

    Iterator it = colListeners.iterator();
    while(it.hasNext())
    {
        ((IEventListener)it.next()).processEvent(mFoo);
    }
}
```

- **Do you see a problem here?**
  - Notification of listeners can be late
  - Member `mFoo` can be changed while updating listeners

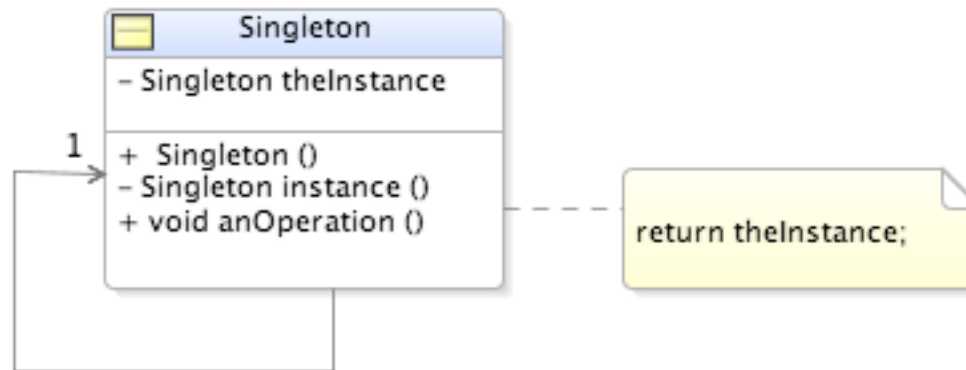
# Assignment & Synchronization

```
public void set(Object foo)
{
    mFoo = foo;

    Iterator it = colListeners.iterator();
    while(it.hasNext())
    {
        ((IEventListener)it.next()).processEvent(foo);
    }
}
```

- **Local variables can resolve some problems**

# Singleton Pattern



- Insures there is one instance of the class
  - Constructor is private



# Singleton enum Implementation

- Use of an enum is the preferred way to implement singletons

```
public enum EnumSingleton {  
    INSTANCE;  
  
    public void anyMethod() {  
        // implementation  
    }  
}
```

# Singleton vs. Class Methods

- **Benefits over a class having only class (static) methods**
  - Constructors apply to instances
  - Only instances can be serialized
  - Interfaces can't specify static methods
  - Inheritance is not useful
  - May insert an object pool if a single instance is not sufficient (non-enum implementation)

# Lazy Initialization Class Holder Idiom

- Lazily initializes a static field
  - Classes are guaranteed not to be initialized until used
  - Useful for singleton implementation pre 1.5

```
public class MySingleton {  
    private static class InstanceHolder {  
        static final MySingleton singleton = new MySingleton();  
    }  
  
    private MySingleton() {  
    }  
  
    public static MySingleton instance() {  
        return InstanceHolder.singleton;  
    }  
}
```

# Lazy Initialization Double-check Idiom

- Double-check idiom
  - Broke prior to 1.5, due to weak `volatile` semantics

```
public class LazyInitialization {  
    private volatile LazyType lazyValue;  
  
    public LazyType getLazyValue() {  
        LazyType result = lazyValue;  
        if (result == null) {  
            synchronized(this) {  
                if (result == null) {  
                    result = lazyValue = computeInitValue();  
                }  
            }  
        }  
        return result;  
    }  
    ...  
}
```

# Cooperating Threads

# Monitor Methods

- `java.lang.Object`
  - `wait()`
    - Thread must have ownership of object's monitor prior to calling
    - Thread releases ownership, and waits, until...
  - `notify()`
    - Thread must own object's monitor
    - Thread wakes up **any one** of the threads that called `wait()` on the object
    - Awakened thread must obtain ownership of the object's monitor

# Synchronization Example

```
int value = 0;
Object lock = new Object;
void method1()
{
    synchronized(lock)
    {
        ++value;
        try {
            lock.wait();
        }
        catch(InterruptedException ie) {
            // means something called interrupt()
        }
        finally {
            --value;
        }
    }
}
```


# Synchronization Example

```
void method2()  
{  
    synchronized(lock)  
    {  
        lock.notify();  
        --value;  
    }  
}
```

**What is a value of `value` going to be?**



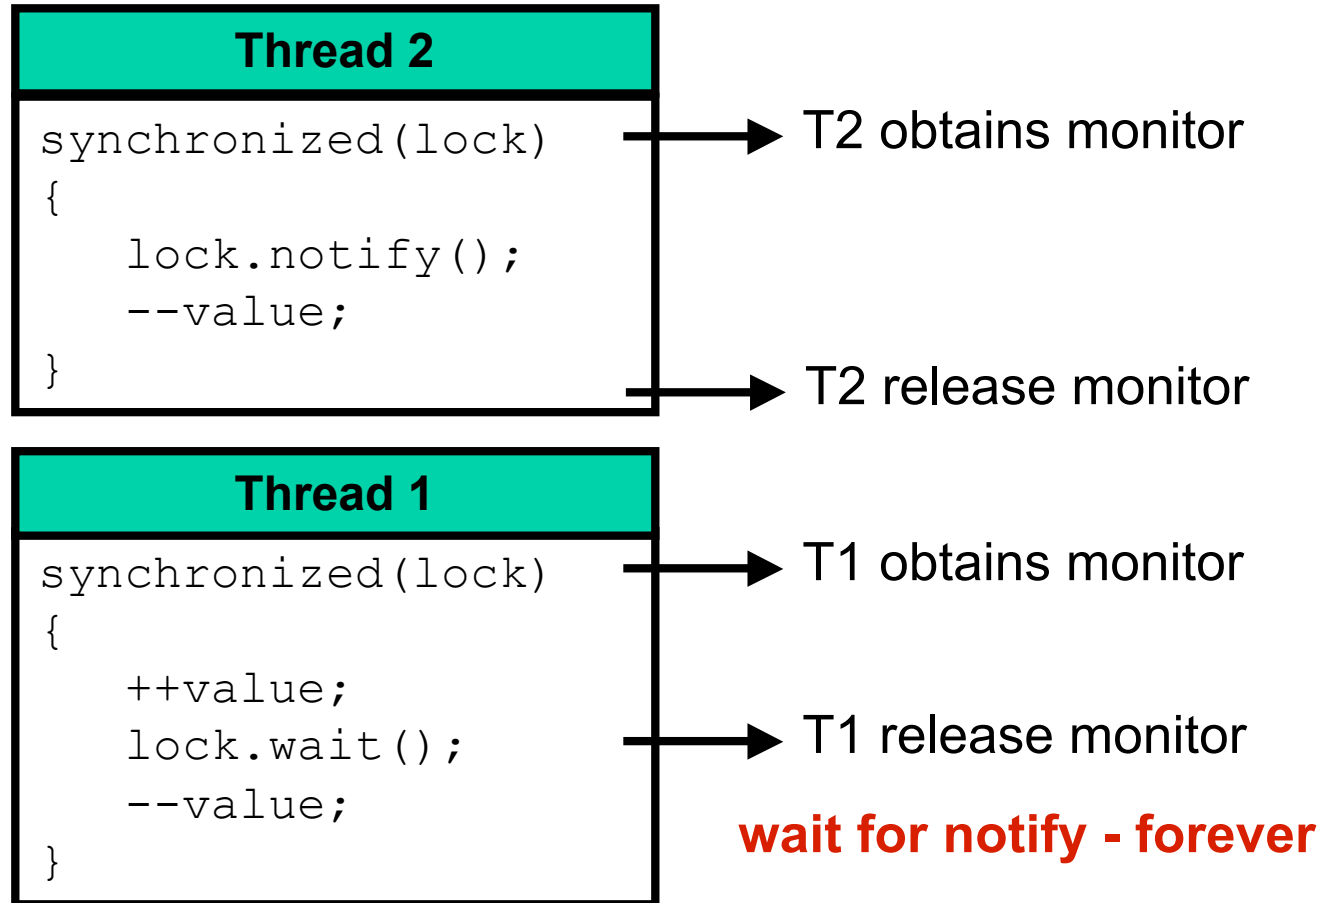
# How it Works



# More Monitor Methods

- `java.lang.Object`
  - `notifyAll()` - **awakens all *wait()*ing threads**
    - **Easy but more expensive**
  - `wait(long timeout)`
    - **Equivalent of `wait()`, but will try to re-obtain the lock and resume execution**
  - `wait(long timeout, int nanos)`
    - **Same as above with “higher precision”**

# Race Condition



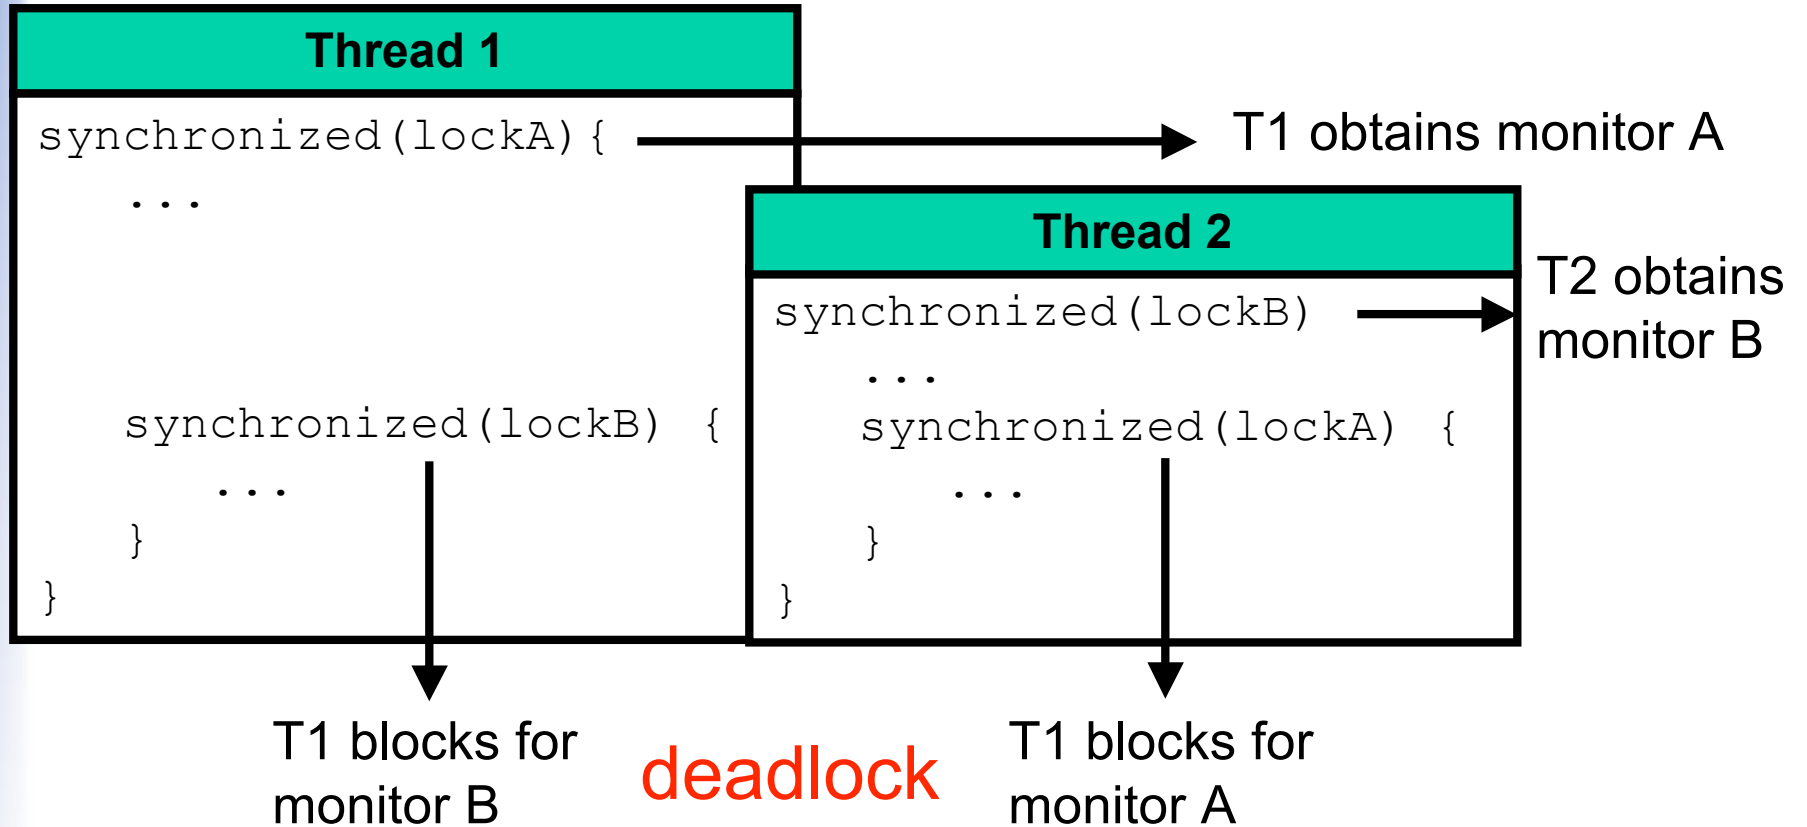
# Race Condition

- Call `wait()` **carefully**, without appropriate `notify()` we have hung thread
- Calling `notifyAll()` can help avoid hung threads
- Only good design will save you! :-)

# Lock Inversion

```
Object lockA = new Object();
Object lockB = new Object();
void method1() {
    synchronized(lockA) {
        ...
        synchronized(lockB) {
            ...
        }
    }
}
void method1() {
    synchronized(lockB) {
        ...
        synchronized(lockB) { {
            ...
        }
    }
}
```

# Lock Inversion - Deadlock



# Coordinating Worker Threads

## Producer - Consumer

- Both threads work on a specific task
  - Producer thread populates queue
  - Consumer picks the results

```
// consumer thread
public void run()
{
    while( true )
    {
        synchronized(list)
        {
            while(list.isEmpty())
                list.wait()
            Object task = list.getFirst();
        }
    }
}
```

# Coordinating Worker Threads

## Producer - Consumer

- **Producer thread**
  - add
  - notify

```
// producer thread
public void run()
{
    synchronized(list)
    {
        list.addLast(task);
        list.notify();
    }
}
```



# Coordinating Worker Threads

## AND-style

- **Workers need to coordinate their work**
  - Both tasks must complete prior to proceeding
  - **AND-style joiner**

```
Thread th1 = new Thread(runnable1);
Thread th2 = new Thread(runnable2);

th1.start();
th2.start();
// both threads running

// wait for first thread to terminate
th1.join();

// wait for second thread to terminate
th2.join();
```

# Coordinating Worker Threads

## OR-style

- Using callback
  - Each thread knows about others
  - When thread completes notifies other threads
  - Notified threads terminate
  - All joins cease blocking
- Use same `join()` scheme

# Shutting Down a Program

- To clean up and shut down a running Java program from the outside (from the command line or a script file)
- Use a *shutdownHook*
  - Simply an initialized but unstarted thread
  - **Must extend** Thread, Runnable doesn't work here
  - All shutdownHooks will be run when the Java runtime is told to stop; by a Cntl-C from the command line, for example

# Shutting Down a Program (cont'd)

- **Instantiate the shutdownHook**
- **Get the current runtime object by calling**  
`Runtime.getRuntime();`
- **Register the shutdownHook by calling**  
`runtime.addShutdownHook(shutdownHook);`
- **When the runtime shuts down, all the registered shutdownHooks will run**